



Pacific58

Jean Paul

November 11, 2024

Contents

1	Introduction	3
	Why Alpine?	3
	Security-Focused	3
	Versatility and Adaptability	3
	Performance and Speed	4
	POSIX Compliant Features of Alpine Linux	4
2	Features	6
	Alpine Package Keeper	6
	APK Metadata	6
3	Deployment	10
	Prerequisites	10
	Deployment Script	10
	Note for Windows Users	11
	Removal Script	11
4	Configurations	13
	Networking	13
	Networking Files	13
	Scripts in Action	13
	Scenario Wrap Up	17
5	Security	18
	PaX	18
	Position Independent Executables (PIE)	21
	Address Space Layout Randomization	21
	Stack Smashing Protection (SSP)	23
	Stack Canaries	24
	Control Flow Integrity (CFI)	27
	Data Execution Prevention (DEP)	29
	Non-Executable Stacks	31
	Kernal Hardening	32

1 Introduction

In today's fast-paced technological landscape, the choice of a robust and efficient operating system is paramount for maintaining a competitive edge. Alpine Linux has emerged as a frontrunner in the domains of IoT, containers, and embedded systems, and for good reason. This section delves into the myriad advantages that make Alpine Linux the ideal choice for our container environment.

Why Alpine?

- ➔ **Footprint:** Alpine Linux is renowned for its minimal footprint.
- ➔ **Size:** Base image size of approximately 5 MB.
- ➔ **Lightweight Solution:** Provides an exceptionally lightweight solution compared to other distributions.
- ➔ **Overhead Reduction:** Particularly advantageous for containerized applications, where reducing overhead is crucial.
- ➔ **Utilities:** Utilizes musl libc and busybox utilities for a compact size and streamlined functionality.
- ➔ **Efficiency:** Ensures that container environments are both fast and resource-efficient.
- ➔ **Why not Windows?:** While Windows is a robust and user-friendly operating system, it is not designed with containers in mind. Windows containers are larger, require more resources, and have more overhead compared to lightweight Linux containers. Additionally, licensing costs and the complexity of managing updates can be a deterrent for using Windows in a container environment.

Security-Focused

- ➔ **♥ PaX:** Alpine Linux employs PaX features for enhanced protection against security threats.
- ➔ **Secure Design:** The security-oriented design ensures robust defense against potential vulnerabilities.
- ➔ **Peace of Mind:** Provides peace of mind by reducing the risk of breaches and ensuring the safety of our container environments.

Versatility and Adaptability

- ➔ **Wide Range of Applications:** Alpine Linux can be deployed in various environments, from IoT devices to large-scale container deployments.
- ➔ **Hardware Compatibility:** Compatible with different hardware platforms, including x86, ARM, and more.
- ➔ **Package Repository:** Extensive package repository that supports a wide range of software, making it suitable for diverse deployment scenarios.
- ➔ **Docker Base Image:** Widely used as a base image for Docker containers, providing a lightweight and efficient foundation.
- ➔ **Flexibility:** Adaptable to different use cases, whether for development, testing, or production environments.
- ➔ **Scalability:** Easily scalable, making it a reliable choice for both small-scale projects and large enterprise solutions.
- ➔ **Community Support:** Strong community and active development, ensuring that it stays up-to-date with the latest technologies and best practices.

Performance and Speed

- ➔ **Fast Boot Times:** Alpine Linux's minimalistic design ensures fast boot times, which is crucial for quickly deploying and scaling containerized applications.
- ➔ **Resource Efficiency:** The lightweight nature of Alpine Linux means it consumes fewer resources, allowing more efficient use of system resources and improved performance.
- ➔ **Optimized for Containers:** Alpine Linux is specifically optimized for container environments, providing a streamlined experience with minimal overhead.
- ➔ **Reduced Startup Times:** Containers based on Alpine Linux have reduced startup times, enabling rapid deployment and scaling of applications.
- ➔ **Improved Application Performance:** The use of musl libc and busybox utilities contributes to better performance and faster execution of applications.
- ➔ **Smaller Footprint:** The small base image size of Alpine Linux means less storage space is required, leading to faster pull times and reduced disk usage.
- ➔ **Network Efficiency:** With smaller image sizes and efficient design, Alpine Linux-based containers require less bandwidth for pulling and deploying images, leading to improved network efficiency.

POSIX Compliant Features of Alpine Linux

- ➔ **POSIX Compliance:** Alpine Linux adheres to the POSIX (Portable Operating System Interface) standards, ensuring compatibility with other Unix-like systems.
- ➔ **musl libc:** Alpine uses musl libc, a lightweight and efficient implementation of the standard C library, which is fully POSIX-compliant. This ensures consistency and reliability across applications.
- ➔ **BusyBox:** The use of BusyBox in Alpine Linux provides a single binary with many standard Unix utilities, all designed to be POSIX-compliant. This helps in maintaining a minimal and efficient system while ensuring compatibility.
- ➔ **Shell Scripting:** The default shell in Alpine Linux is Ash (Almquist Shell), a lightweight POSIX-compliant shell. This allows for the execution of POSIX-compliant shell scripts, ensuring portability and compatibility across different environments.
- ➔ **Core Utilities:** Alpine Linux includes a set of core utilities that follow POSIX standards, ensuring consistent behavior across different systems and applications.
- ➔ **Interoperability:** POSIX compliance in Alpine Linux enhances interoperability with other Unix-like systems, making it easier to port applications and scripts between environments.
- ➔ **Consistency:** By adhering to POSIX standards, Alpine Linux ensures a consistent and predictable environment, which is crucial for developing and deploying reliable applications.
- ➔ **Portable Applications:** POSIX compliance allows applications developed on Alpine Linux to be easily ported to other POSIX-compliant systems, reducing development and maintenance efforts.

2 Features

Alpine Package Keeper

The Alpine Package Keeper, commonly referred to as **apk**, is the package management system used by Alpine Linux. It is designed to be lightweight, fast, and efficient, making it well-suited for resource-constrained environments like containers and embedded systems.

Key Features of apk

- ➔ **Lightweight:** apk is designed with minimalism in mind, ensuring a small footprint while providing powerful package management capabilities.
- ➔ **Fast and Efficient:** apk operates quickly, making package installation, updates, and removal tasks swift and efficient.
- ➔ **Dependencies Management:** Handles package dependencies automatically, ensuring that all necessary dependencies are installed, updated, or removed as needed.
- ➔ **Security:** Supports digital signatures for packages, ensuring the integrity and authenticity of the software.
- ➔ **Repositories:** Provides access to a wide range of software repositories, allowing users to install a diverse set of packages.
- ➔ **Simplicity:** apk commands are straightforward and easy to use, making it accessible even for those new to Alpine Linux.

APK Metadata

when caching is enabled, each repository in Alpine Linux creates or appends to a gzipped tar file within the `/var/cache/apk/` directory. This file contains metadata about the packages available in the repository. Your repositories are specified within the `/etc/apk/repositories` file. Let's break down its contents and structure.

^ Alpine Package Keeper

^ APK Metadata

- ➔ **C**: This field represents the checksum of the package.
- ➔ **P**: The package name.
- ➔ **V**: Version of the package, including the release number
- ➔ **A**: Architecture for which the package is built (e.g., x86_64, armhf)
- ➔ **S**: Package size in bytes
- ➔ **I**: Installed size in bytes
- ➔ **T**: Short description of the package
- ➔ **U**: URL for the package's source or project page
- ➔ **L**: License under which the package is distributed
- ➔ **o**: Origin, which is usually the base package name
- ➔ **m**: Maintainer's name and email address
- ➔ **t**: Build timestamp in Unix epoch format
- ➔ **c**: Package's commit or checksum identifier
- ➔ **D**: Dependencies required by the package
- ➔ **p**: Provides information on which packages this package can provide or replace
- ➔ **k**: Whether the package is a subpackage
- ➔ **M**: Message shown post-installation or post-removal
- ➔ **i**: information about installed files
- ➔ **e**: Explanation of ABI compatibility

Script for Extracting and Displaying APKINDEX Metadata

Due to length constraints, the full script for extracting and displaying metadata from the APKINDEX files is hosted on GitHub. You can find the complete script at the following link: [View the Full Script on GitHub](#)

^ Alpine Package Keeper

Package Installation

- ➔ **-no-cache**: The option fetches the latest package information from the repositories and does not store the package index locally.
- ➔ **-update-cache** or **-U**: This flag forces a cache update before installing the package.
- ➔ **-virtual** or **-t**: This flag allows you to create a virtual package that groups multiple packages together.
- ➔ **-interactive** or **-i**: This flag prompts for confirmation before installing packages.
- ➔ **-no-scripts** or **-X**: This flag skips running trigger scripts during package installation.
- ➔ **-repository** or **-X**: This flag specifies an additional repository to use for the package installation.
- ➔ **-force** or **-f**: This flag forces the installation of a package even if it would replace another package.
- ➔ **-root** or **-p**: This flag installs packages to a different root directory.
- ➔ **-update** or **-U**: This flag updates all the available packages before installing the specified package.
- ➔ **-allow-untrusted** or **-U**: This flag allows the installation of untrusted packages (packages that are not signed with a trusted key).
- ➔ **-wait** or **-w**: This flag waits for the lock to be released before proceeding with the installation.
- ➔ **-clean-protected** or **-P**: This flag cleans the package index after installation, even if it's protected.
- ➔ **-no-progress** or **-q**: This flag suppresses the progress output during installation.
- ➔ **-arch** or **-a**: This flag installs a package for a different architecture.
- ➔ **-simulate** or **-s**: This flag simulates the installation without actually performing it.
- ➔ **-quiet** or **-q**: This flag suppresses all output, making the installation process completely silent.

^ Alpine Package Keeper

1

^ -no-cache

^ Alpine Package Keeper

1

^

^ Alpine Package Keeper

Package Search

- ➔ **-v**: Verbose output
- ➔ **-i**: Search only installed packages
- ➔ **-q**: Quiet output

```
1 # This command will list packages that match man-db in the apk repository
2 apk search man
3
4 # This command will provide a verbose output (-v) of all installed packages
  ↳(-i) that match 'man' in the apk repository.
5 apk -iv search man
6
7 # This command will provide a quiet output (-q) of all installed packages (-i)
  ↳that match 'man' in the apk repository.
8 apk -iq search man
```

Using apk search effectively can help you quickly find the packages you need. These commands combine the **-i** flag to search installed packages, the **-v** flag for verbose output, and the **-q** flag for quiet output, respectively. This ensures that the search results are tailored to your preferred level of detail.

3 Deployment

Prerequisites

Before we begin, ensure you have an up-to-date version of [Docker](#) installed.

Deployment Script

Description: To illustrate the deployment process, below is an example script that demonstrates how to launch an Alpine Linux container with unique identifiers and network configuration. This script generates a unique prefix for the container, runs it in detached mode with the host's network stack, and assigns custom labels, hostname, and domain name. It also keeps the container running persistently, allowing for interactive management via Docker's `exec` command.

Note for Windows Users

- ➔ **Cygwin Installation:** Visit the [Cygwin website](#) to download and install Cygwin.
- ➔ **POSIX Environment:** Cygwin provides a robust POSIX-compatible environment, enabling you to run and manage the examples seamlessly on Windows.



```
1  (  
2      CONTAINER_ID="$(  
3          UNIQUE="$(cat /dev/urandom | tr -dc [:alnum:] | head -c 16)"  
4          docker container run --detach \  
5              --network=bridge \  
6              --name example_${UNIQUE} \  
7              --label ${UNIQUE} \  
8              --hostname ${UNIQUE}.example.lab \  
9              --domainname example.lab \  
10             --publish-all \  
11             alpine:latest ash -c "while ;; do sleep 1; done"  
12     )" && {  
13         docker container exec --interactive --tty ${CONTAINER_ID} ash  
14     }  
15 )
```

Throughout this documentation, we will refer to the provided example script as a foundation for various deployment scenarios and configurations. This script demonstrates best practices for container deployment, including unique naming, network configuration, and persistent container operation. Please ensure you are familiar with this example, as it will serve as a reference point for understanding and implementing the discussed concepts.

Removal Script

Description: This script identifies and removes Docker containers with names matching the pattern `example_XXXXXXXXXX` where `XXXXXXXXXX` is a unique 16-character alphanumeric string. It lists all containers, filters those that match the name pattern, and then stops and removes each matching container. The process is performed in parallel for efficiency.



```
1  (  
2      for I in $(  
3          docker container ls -a --format '{{.Names}} {{.ID}}' | awk '{  
4              if ($1 ~ /^example_[[:alnum:]]{16}$/) {  
5                  print $2  
6              }  
7          }'  
8      ); do  
9          {  
10             docker container stop ${I}  
11             docker container rm ${I}  
12         } &  
13     done  
14 )
```

4 Configurations

Networking

Networking Files

- ➔ ♡ **/etc/network/interfaces**: The main configuration file for network interfaces. It defines how each network interface is set up (e.g., static IP, DHCP).
- ➔ ♡ **/etc/network/interfaces.d**: A directory for additional network interface configurations. Any file in this directory will be included by the main interfaces file. This is useful for organizing configurations separately.
- ➔ ♡ **/etc/network/if-down.d**: Scripts in this directory are executed when an interface is brought down. This can be used for cleaning up or logging.
- ➔ **/etc/network/if-post-down.d**: Similar to if-down.d, but these scripts run after the interface has been fully brought down.
- ➔ **/etc/network/if-post-up.d**: Scripts here are run after an interface is brought up. Useful for post-configuration tasks like setting additional routes or updating DNS.
- ➔ ♡ **/etc/network/if-pre-down.d**: Scripts executed before an interface is brought down. Can be used for pre-shutdown checks or preparations.
- ➔ ♡ **/etc/network/if-pre-up.d**: Scripts that run before an interface is brought up. Useful for preliminary checks or setting up necessary conditions before the interface is up.
- ➔ ♡ **/etc/network/if-up.d**: Scripts executed when an interface is brought up. These can perform tasks such as updating network settings or notifying other services.

Scenario

- ➔ **lo**: The loopback interface.
- ➔ **eth0**: Will be assigned a static IP address.
- ➔ **eth1**: Will be configured to obtain an IP address via DHCP.

^ Networking

^ /etc/network/if-pre-up.d

```
1  #!/bin/ash
2  # /etc/network/if-pre-up.d/rename-interfaces
3  # This script renames network interfaces before they are brought
4  ip link show eth0 1> /dev/null 2>&1 && ip link set dev eth0 name eth0_static
5  ip link show eth1 1> /dev/null 2>&1 && ip link set dev eth1 name eth1_dhcp
```

^ /etc/network/if-up.d

```
1  #!/bin/ash
2  # /etc/network/if-up.d/add-route
3  # This script adds a route after the interface is brought up
4  ip route add 192.168.2.0/24 via 192.168.1.1 dev eth0_static
```

^ /etc/network/interfaces

```
1  #!/bin/ash
2  # /etc/network/interfaces
3  # Main network interfaces configuration file
4
5  # Loopback interface
6  auto lo
7
8  # DHCP configuration for eth1_dhcp
9  auto eth1_dhcp
10 iface eth1_dhcp inet dhcp
```

^ Networking

^ /etc/network/interfaces.d

```
1 #!/bin/ash
2 # /etc/network/interfaces.d/eth0_static
3 # Static IP configuration for eth0_static
4 auto eth0_static
5 iface eth0_static inet static
6     address 192.168.1.100
7     netmask 255.255.255.0
8     gateway 192.168.1.1
9     dns-nameservers 8.8.8.8 8.8.4.4
```

^ /etc/network/if-pre-down.d

```
1 #!/bin/ash
2 # /etc/network/if-pre-down.d/pre-down-check
3 # Pre-down check script
4 # Perform necessary actions before the interface is brought down
5
6 # Example: Logging the interface down event
7 logger "Bringing down interface: ${IFACE}"
```

^ /etc/network/if-down.d

```
1 #!/bin/ash
2 # /etc/network/if-down.d/cleanup-routes
3 ip route del 192.168.2.0/24 via 192.168.1.1 dev eth0_static
```

^ Networking

Alpine Linux VM: Use this setup if you are working with a full virtual machine (VM) environment. This is ideal for scenarios requiring a complete OS with persistent storage, system-level configurations, and traditional init systems like OpenRC.



```
1 # Restart the networking service to apply the configurations
2 service networking restart
3
4 # Check the status of the networking service
5 service networking status
```

^ Docker Container: Use this setup if you are running lightweight, isolated applications in containers. This approach is perfect for microservices, testing, and deployment scenarios where quick startup times and efficient resource usage are critical.



```
1 # Inside the container: Install necessary packages
2 apk update && apk upgrade
3 apk add openrc iproute2
4
5 # Create the /run/openrc directory if it doesn't exist and create the
↳softlevel file to bypass OpenRC warnings
6 # This is required because Docker containers do not typically have an init
↳system like OpenRC by default.
7 # Without this, you may encounter warnings when trying to manage services with
↳OpenRC.
8 mkdir -p /run/openrc && touch /run/openrc/softlevel
9
10 # Inside the container: Enable and start the networking service
11 rc-update add networking default
12 rc-service networking start
13
14 # Inside the container: Verify the networking service status
15 rc-service networking status
```


^ Networking

Conclusion: In this section, we delved into the networking configurations and scripts of Alpine Linux. Here's a summary of what we covered:

Scenario Wrap Up

- ➔ **Interface Setup:** Configured lo (the loopback interface), eth0 with a static IP address, and eth1 to obtain an IP address via DHCP.
- ➔ **Scripts:** Implemented scripts to rename interfaces before they are brought up and add routes after the interfaces are up.
- ➔ **Cleanup:** Provided cleanup scripts for routes before interfaces are brought down.

What We Did

- ➔ **IP Configurations:** Demonstrated configurations for static and dynamic IP addresses.
- ➔ **Script Examples:** Included example scripts to handle network interface changes, ensuring proper setup, cleanup, and logging.

Wrap Up

By following these examples, you should have a solid foundation for managing network interfaces in Alpine Linux. This structured approach helps keep your network setup organized and efficient.

5 Security

PaX

Description: PaX provides a set of security enhancements originally derived from the PaX patches for the Linux kernel. These include various protections against memory corruption vulnerabilities.

Pax Flags

- ➔ **▼ PAGEEXEC:** Prevents execution of code on writable pages (stack and heap). This is one of the key protections against code injection attacks, ensuring that even if an attacker manages to inject malicious code, it cannot be executed.
- ➔ **▼ EMUTRAMP:** Emulates trampolines to allow certain legitimate uses of writable and executable pages. This flag makes exceptions for specific, safe cases where executing writable pages is necessary, like for some just-in-time (JIT) compilers.
- ➔ **▼ MPROTECT:** Forces the use of mprotect to set memory protections, ensuring that memory protections are enforced. By restricting permissions changes to memory pages, it prevents common exploits that rely on making a page writable and executable.
- ➔ **▼ RANDMMAP:** Implements Address Space Layout Randomization (ASLR) to randomize the memory layout of processes, making it harder for attackers to predict addresses. This increases the complexity of attacks by making it difficult for attackers to predict the location of specific functions or variables in memory.
- ➔ **▼ SEGMEXEC:** Allows the execution of code in the data segment, providing an additional layer of protection against buffer overflow attacks. This protection creates separate memory segments for executable and non-executable data, reducing the risk of executing malicious code.

^ PaX

PaX Setup

```
1 # Ensure pax-utils and paxctl are installed in the container
2 apk update && apk upgrade
3 apk add pax-utils paxctl
4
5 # Convert the PT_GNU_STACK program header to PT_PAX_FLAGS for busybox
6 paxctl -c /bin/busybox
7
8 # Check the current PaX flags for the busybox executable
9 paxctl -v /bin/busybox
```

^ PAGEEXEC

```
1 # Enable PAGEEXEC for an executable
2 # This ensures that code cannot be executed from writable pages
3 paxctl -P busybox
4
5 # Disable PAGEEXEC for an executable
6 # This allows execution of code from writable pages (not recommended for
↳security)
7 paxctl -p busybox
```

^ EMUTRAMP

```
1 # Enable EMUTRAMP for an executable
2 # This allows certain legitimate uses of writable and executable pages
3 paxctl -E busybox
4
5 # Disable EMUTRAMP for an executable
6 # This disables the emulation of trampolines
7 paxctl -e busybox
```

^ PaX

^ MPROTECT

```
1 # Enable MPROTECT for an executable
2 # This forces the use of mprotect to set memory protections
3 paxctl -M busybox
4
5 # Disable MPROTECT for an executable
6 # This disables the enforcement of memory protections
7 paxctl -m busybox
```

^ RANDMMAP

```
1 # Enable RANDMMAP for an executable
2 # This implements ASLR to randomize the memory layout of processes
3 paxctl -R busybox
4
5 # Disable RANDMMAP for an executable
6 # This disables the randomization of memory addresses
7 paxctl -r busybox
```

^ RANDEXEC

```
1 # Enable RANDEXEC for an executable
2 # This randomizes the locations of executable segments
3 paxctl -X busybox
4
5 # Disable RANDEXEC for an executable
6 # This disables the randomization of executable segments
7 paxctl -x busybox
```

^ SEGMEXEC

```
1 # Enable SEGMEXEC for an executable
2 # This allows execution of code in the data segment
3 paxctl -S busybox
4
5 # Disable SEGMEXEC for an executable
6 # This disables execution of code in the data segment
7 paxctl -s busybox
```

^ PaX

Wrap Up

By integrating PaX, Stack Smashing Protection, and Position Independent Executables, Alpine Linux ensures a robust security posture against various types of attacks. These features help to mitigate the risks of memory corruption, buffer overflows, and code injection attacks, providing a secure environment for running applications.

Position Independent Executables (PIE)

Description: Position Independent Executables (PIE) are binaries designed to be loaded at any memory address without modification. This flexibility allows the operating system to load them at different addresses each time they are executed, providing significant security benefits.

Address Space Layout Randomization (ASLR)

- ➔ **Security:** By randomizing the memory addresses used by executables, PIE makes it significantly more difficult for attackers to predict where specific functions or data are located. This randomness helps protect against various types of exploits, such as buffer overflow attacks.
- ➔ **Implementation:** When a PIE binary is compiled, it does not assume it will be loaded at any fixed address, allowing the system's loader to place it anywhere in the virtual memory space.

Flexibility

- ➔ **Compatibility:** PIE binaries can be shared across different operating systems and architectures without requiring specific address space allocations.
- ➔ **Optimization:** Although there may be slight performance overheads due to the additional relocations needed during execution, the security benefits usually outweigh these costs.

^ Position Independent Executables (PIE)

Dynamic Linking

- ➔ **Efficiency:** PIE works well with dynamically linked libraries (DLLs), as both can be loaded at random addresses, further enhancing security.
- ➔ **Functionality:** By enabling shared code, dynamic linking helps save memory and allows for modular updates and patches, improving system maintenance.

Demonstration

```
1 apk update && apk upgrade
2 apk add gcc build-base
3 cd /tmp
4
5 {
6 cat << 'EOF'
7 #include<stdio.h>
8
9 int main() {
10     printf("Hello, World!\n");
11     return 0;
12 }
13 EOF
14 } > example.c
15
16 # Compile a program as a PIE explicitly
17 gcc -pie -fPIE -o tmpa example.c
18
19 # Compile a program without explicit PIE flags
20 gcc -o tmpb example.c
21
22 # Verify if the binary is position-independent
23 readelf -h tmpa | grep Type | awk '{print $2}' # Should show "DYN"
24 readelf -h tmpb | grep Type | awk '{print $2}' # Should also show "DYN"
```

Note

The two commands produce the same result because, by default, Alpine Linux compiles executables as position-independent. When you compile a program using **gcc** on Alpine Linux, it generates position-independent code by default, due to the way the toolchain is configured.

^ Position Independent Executables (PIE)

Flags

- ➔ **-pie**: This flag is used by the linker to create an executable that supports being loaded at different memory addresses.
- ➔ **-fPIE**: This flag instructs the compiler to generate position-independent code. This means that the code does not assume any specific memory addresses for functions or variables.
- ➔ **-pie -fPIE**: When used together, **-fPIE** ensures that the code generated is position-independent, and **-pie** ensures that the resulting executable can be loaded at any memory address.

Key Points

- ➔ **Default PIE Compilation**: Alpine Linux's GCC toolchain is configured to create Position Independent Executables (PIEs) by default. This means that even without specifying **-pie** and **-fPIE**, the resulting binaries are position-independent.
- ➔ **Same Binary Type**: When you run `readelf -h` on both binaries, they both show `Type: DYN`, indicating that they are dynamically linked and position-independent.

Wrap Up

Position Independent Executables (PIE) enhance the security and flexibility of binaries by allowing them to be loaded at random memory addresses, making it much harder for attackers to exploit vulnerabilities. By incorporating PIE, systems can achieve a higher level of security through techniques like ASLR, dynamic linking, and more.

Stack Smashing Protection (SSP)

Description: what is stack smashing? Stack smashing, also known as a stack buffer overflow, is a type of security vulnerability that occurs when a program writes more data to a buffer (a temporary storage area) on the stack than it can hold. This excess data overwrites adjacent memory, which can corrupt the stack and potentially allow an attacker to execute arbitrary code.

Breakdown

- ➔ **Buffer Overflow:** The program writes data beyond the allocated buffer space, spilling over into adjacent memory regions on the stack.
- ➔ **Stack Corruption:** The overflowed data may overwrite important control information, such as the return address of a function.
- ➔ **Exploitation:** An attacker can exploit this by carefully crafting the overflow data to overwrite the return address with the address of malicious code, causing the program to execute the attacker's code when the function returns.

Stack smashing attacks can lead to serious security breaches, including unauthorized access and system crashes. To mitigate these risks, modern compilers and operating systems implement various protections, such as [stack canaries](#), [ASLR](#), and [non-executable stacks](#).

Alpine Linux compiles its binaries with stack smashing protection out of the box. This typically involves adding stack canaries, which are special values placed on the stack that get checked before a function returns. If the canary value has been altered, the program will detect a buffer overflow and exit gracefully instead of continuing to execute potentially malicious code.

Stack Canaries

Description: Stack canaries are a security mechanism used to detect and prevent stack buffer overflow attacks, which occur when a program writes more data to a buffer on the stack than it can hold.

Breakdown

- ➔ **Canary Value:** A small, known value (the "canary") is placed between the stack's control data (such as the return address) and the local variables.
- ➔ **Checking the Canary:** Before a function returns, the canary value is checked to see if it has been altered.
- ➔ **Detection of Overflow:** If a buffer overflow occurs and overwrites the canary value, the program detects the change and typically responds by terminating or taking other protective actions.

The goal of stack canaries is to prevent the attacker from overwriting important control data, such as return addresses, without being detected.

Types

- ➔ **Terminator Canaries:** These canaries contain null bytes, carriage returns, and newline characters, making them difficult to overwrite with standard string functions.
- ➔ **Random Canaries:** These canaries are randomly generated at program startup, making them unpredictable and harder for an attacker to guess.
- ➔ **Guard Pages:** These canaries use memory pages marked as non-readable or non-writable to detect stack overflows.

Stack canaries are an important part of modern software security practices, helping to protect against certain types of buffer overflow attacks.

The script below demonstrates how to create and run two C programs—one with stack canaries enabled and one without. The programs are designed to illustrate how stack canaries can detect and prevent buffer overflow attacks.

^ Stack Canaries

^ Example: Terminator Canary

```
1 apk update && apk upgrade
2 apk add clang build-base
3 cd /tmp
4
5 {
6 cat << 'EOF'
7 #include <stdio.h>
8 #include <string.h>
9
10 void vulnerable_function(char*);
11
12 int main(int argc, char **argv) {
13     if (argc != 2) {
14         fprintf(stderr, "Usage: %s <input string>\n", argv[0]);
15         return 1;
16     }
17
18     vulnerable_function(argv[1]);
19     printf("Function executed successfully.\n");
20     return 0;
21 }
22
23 void vulnerable_function(char *str) {
24     char buffer[16];
25     strcpy(buffer, str); // This is intentionally vulnerable
26 }
27 EOF
28 } > vulnerable_example.c
29
30 # Compile the file without stack canaries
31 gcc -fno-stack-protector -o example_without_canary vulnerable_example.c
32
33 # Compile the file with stack canaries enabled (default)
34 gcc -o example_with_canary vulnerable_example.c
35
36 # Run the program with stack canaries
37 # (this should detect the overflow)
38 ./example_with_canary "this_is_a_very_long_input_string"
39
40 # Run the program without stack canaries
41 # (this might crash or behave unexpectedly)
42 ./example_without_canary "this_is_a_very_long_input_string"
```

^ Stack Canaries

Note

Since this is Alpine, you should expect the program without stack canaries to result in a core dump when a buffer overflow occurs. This demonstrates the robustness of Alpine's handling of such vulnerabilities.

Wrap Up

We provided examples of simple C programs, one with stack canaries enabled and one without, to demonstrate how canaries help detect and prevent buffer overflow attacks.

Stack canaries play a crucial role in modern software security by providing a line of defense against stack buffer overflow attacks, which can lead to serious security breaches, unauthorized access, and system crashes.

Stack canaries are just one of many techniques used to enhance software security, but their effectiveness in detecting and preventing buffer overflows makes them a valuable tool in the fight against exploitation.

Control Flow Integrity (CFI)

Description: Control Flow Integrity (CFI) is a security technique designed to ensure that the control flow of a program remains unaltered during execution. It prevents attackers from hijacking the program's control flow, which is a common goal in many exploits, such as buffer overflow attacks. By enforcing CFI, the program can only execute legitimate control flow paths, making it much harder for attackers to execute arbitrary code.

How CFI Works

- ➔ **Instrumenting Code:** Adding checks to the program's code to ensure that indirect jumps and calls only go to valid targets.
- ➔ **Creating Control Flow Graphs (CFG):** Mapping out all possible paths the program can take during execution.
- ➔ **Checking Transfers:** At runtime, verifying that control transfers (like function calls and returns) follow the paths defined in the CFG.

^ Control Flow Integrity (CFI)

CFI Example Script

```
1 apk update && apk upgrade
2 apk add clang build-base
3 cd /tmp
4
5 {
6 cat << 'EOF'
7 #include <stdio.h>
8 #include <stdlib.h>
9 #include <stdint.h>
10
11 void target_function() {
12     printf("Target function executed.\n");
13 }
14
15 void vulnerable_function(void (*func)()) {
16     func();
17 }
18
19 int main(int argc, char **argv) {
20     if (argc != 2) {
21         fprintf(stderr, "Usage: %s <input>\n", argv[0]);
22         return 1;
23     }
24
25     void (*func)() = NULL;
26
27     if (argv[1][0] == '1')
28         func = target_function;
29     else
30         func = (void (*)())(uintptr_t)strtoul(argv[1], NULL, 0);
31
32     vulnerable_function(func);
33
34     return 0;
35 }
36 EOF
37 } > example.c
38
39 TMPA="$(clang -fsanitize=cfi \
40     -flto -fvisibility=hidden \
41     -o cfi_demo example.c 2>&1 | awk -F "''" '{print $2}')"
42
43 mkdir -p $(dirname "${TMPA}")
44 # Continues on the next page...
```

^ Control Flow Integrity (CFI)

^ CFI Example Script

```
1 {
2 cat << 'EOF'
3 fun:func
4 EOF
5 } > "${TMPA}"
6
7 unset TMPA
8
9 clang -fsanitize=cfi -flto \
10     -fvisibility=hidden \
11     -o cfi_demo example.c
12
13 # Valid input (target_function will be executed)
14 ./cfi_demo 1 # Output: Target function executed.
15
16 # Invalid input (CFI should detect control flow hijacking)
17 ./cfi_demo 0x12345678
18
19 objdump -d cfi_demo | grep '.cfi'
```

Note

While GCC has some support for CFI, Clang's implementation of CFI is more mature and reliable. Clang provides better instrumentation and runtime checks, making it a preferred choice for enabling advanced security features like CFI. This is why we used Clang in our demonstration to ensure robust and effective protection against control flow hijacking attacks.

Wrap Up

We also demonstrated how to compile a C program with CFI enabled using Clang, and verified the presence of CFI instrumentation in the binary using `objdump`.

Data Execution Prevention (DEP)

Description: DEP is a security feature that helps prevent code from being executed in certain regions of memory that are not intended to contain executable code. This can help protect against attacks that exploit vulnerabilities by injecting malicious code into a program's data areas, such as the stack or heap.

How DEP Works

DEP marks specific areas of memory as non-executable. If a program attempts to run code from these non-executable regions, DEP intervenes and prevents the execution, often resulting in the program being terminated. This helps to mitigate various types of attacks, such as buffer overflow and stack-based attacks, where malicious code is injected into data areas.

Types of DEP

- ➔ **Hardware-enforced DEP:** Utilizes hardware capabilities in modern CPUs to mark memory regions as non-executable. This is the most effective form of DEP.
- ➔ **Software-enforced DEP:** Provides similar protection through software mechanisms. It is used in environments where hardware DEP is not available or supported.

Benefits of DEP

- ➔ **Preventing malicious code execution:** Prevents malicious code from executing in non-executable memory regions.
- ➔ **Reducing attack surface:** Reduces the attack surface for various exploits, enhancing overall system security.
- ➔ **Integration with other security measures:** Works in conjunction with other security measures, such as ASLR and CFI, to provide comprehensive protection.

^ Data Execution Prevention (DEP)

DEP Example

```
1 apk update && apk upgrade
2 apk add gcc build-base
3 cd /tmp
4
5 {
6 cat << 'EOF'
7 #include <stdio.h>
8 #include <string.h>
9
10 void target_function();
11 void vulnerable_function(char*);
12
13 int main(int argc, char **argv) {
14     if (argc != 2) {
15         fprintf(stderr, "Usage: %s <input string>\n", argv[0]);
16         return 1;
17     }
18
19     vulnerable_function(argv[1]);
20     printf("Function executed successfully.\n");
21     return 0;
22 }
23
24 void target_function() {
25     printf("Target function executed.\n");
26 }
27
28 void vulnerable_function(char *str) {
29     char buffer[16];
30
31     // This is intentionally vulnerable
32     strcpy(buffer, str);
33 }
34 } > example.c
35 gcc -z noexecstack -fno-stack-protector -o dep_demo example.c
36 ./dep_demo "short_input"
37 ./dep_demo "this_is_a_very_long_input_string_that_overflows_the_buffer"
38
```

Non-Executable Stacks

Description:

Kernal Hardening

Description: